

Jai Dhyani

Last Updated on 1st May 2026
github.com/jaidhyani/resume

jai.one | github.com/jaidhyani | jai@jai.one | it's pronounced "J"

I want the future to be full of human flourishing.

To that end I build things and bash them against reality to figure out why they don't work, and then I try to build something better.

Co-author on RE-Bench (METR's AI R&D time-horizon eval). Founder, Luthien Research.

Previously ML for hate speech at Facebook (the intersection of human values and AI is very hard).

LUTHIEN RESEARCH luthien.cc | github.com/luthienResearch/luthien-proxy | FOUNDER AND CTO

2025 - 2026: Technical AI Safety Nonprofit | 2026 - Present: Public Benefit Corporation

- Founded a startup to implement AI control in live production systems at the API level
- Built a free, open-source LLM API proxy; first to provide API-level tool call edits + insertions, multi-session context, and live modification of streaming responses
- Deployed at Redwood Research and Trajectory Labs
- Raised over \$350k

MODEL EVALUATION AND THREAT RESEARCH (METR) | VISITING FELLOWSHIP

2024 via ML Alignment and Theory Scholars (MATS) 6.0

- Co-author on "RE-Bench: Evaluating frontier AI R&D capabilities of language model agents against human experts"
- Developed evaluations and tooling to detect AI research capabilities as part of Hjalmar Wijk's team

INDEPENDENT PROJECTS

2023: Delphi - Evals Technical Lead

- Led team building infrastructure for training and evaluating small, reproducible language models
- Rewrote training pipeline for improved reliability, speed, and configurability

2023 - 2024: Rational Animations Scriptwriter

TEMPLE CAPITAL | ML ENGINEERING CONSULTANT

2022

- Sped up quantitative investment analysis engine; added CI/CD pipeline
- Extended trading system to handle synthetic assets

BLUE ROSE RESEARCH | SENIOR ML ENGINEER

2021

- Accelerated TF-Probability variational inference pipelines for public opinion modeling
- Diagnosed and fixed slow database queries in analytical workflows

META (FACEBOOK) | ML ENGINEERING FOR HATE SPEECH

2017 - 2020: Software Engineer - Machine Learning & Infrastructure

- Built Facebook's first proactive hate speech classifiers; integrated CV and multilingual NLP
- Built volume-adjusted thresholding system to optimize experts' time and maximize per-sample training value
- Discovered novel failure modes in training infrastructure; developed stratified offline datasets to prevent regression
- Trained and deployed a model that deleted the Declaration of Independence from Facebook on July 4th 2018, subsequently redoubling my respect for Goodhart's Law.

LENDINGROBOT | SENIOR BACKEND ENGINEER

2016 - 2017

AMAZON | SYSTEMS ENGINEER II, AWS NETWORKING

2014 - 2015

DEMOCRATIC NATIONAL COMMITTEE | LINUX SYSTEMS ENGINEERING

2011 - 2012: Linux Systems Engineer | 2013: Lead Linux Systems Engineer

EDUCATION

University of Chicago | B.S. Computer Science | 2011